

notebook test and validation metrics of the various models are shown and compared. A selection of the best performing model will end this project.

```
: reset -fs

: # Importing the most important modules and setting the style for following plots
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from matplotlib.pyplot import cm
import warnings
from datetime import datetime
import pickle
from scipy.integrate import simps

import plotly.offline as pyo
import plotly.graph_objects as go

# For Data Mining and data export
import glob
```

```
# Setting the random seed for reproducability and several plotting style parameters
%matplotlib inline
warnings.filterwarnings('ignore')
pd.set_option('display.max_columns', None)

# Set notebook mode of plotly to work offline
pyo.init_notebook_mode()

: def split_at(string, separator, at_n, return_after=True):
    """
    helper function to sperate file path into name
    """
    words = string.split(separator)
    if return_after:
        return separator.join(words[at_n:])
    else:
        return separator.join(words[:at_n])

def clean_brackets(df_result):
    """
    Explanation: Working with LSTMs, the error metrics are sometimes encapsulated into square brackets.

    Return: dataframe without square brackets.
    """
    for colm in df_result.columns:
        if df_result[colm].dtype != np.number:
            df_result[colm] = [float(x.replace("[", "").replace("]", "")) for x in df_result[colm]]
    return df_result

: def get_errros(metric = 'MAPE', folder='Test Errors'):
    """
    Reason: Error Metrics are saved into .csv files, one file per model including multiple error metrics

    metric: ['RMSE', 'R2', 'MAE', 'MAPE']
    folder: ['Test Errors', 'Validation Errors']
    Return: pd.DataFrame with specified Error Metric for all .csvs in specified folder
    """
    file_start = '_'.join(folder.split(' ')).lower()
    errors = pd.DataFrame()
    for file in glob.glob(f'{folder}/{file_start}*.{metric}'):
        # Preparation of column Name
        name = file
        name = split_at(file, '_', 2, return_after=True)
        name = split_at(name, '.', 1, return_after=False)
        name = name.replace('_', '_').replace('.', '.')
```

```
    raw_table.index +=  
    errors = pd.concat([  
        errors.index.rename('S'),  
        errors], axis=1)  
    errors = clean_brackets(errors)  
    errors = errors.sort_index()  
  
    return errors
```

CO
ta'] #
if

```

plt.plot(df_results[colm], label=colm ,color=c, lw=4)

##### PLOT SETTINGS #####
plt.title(f'{metric} for each predicted timestep into the future ({val_or_}
plt.xlabel('STEPS INTO THE FUTURE', fontsize=14)
plt.ylabel(f'{metric}', fontsize=14)
plt.xlim(1, 19)
plt.legend(loc='upper left', fontsize=14, bbox_to_anchor=(1.01, 1))

##### OUTPUT #####
# plt.savefig(f'figures/results_{metric}_{val_or_test}.png', bbox_inches='tight')
plt.show();

: def plot_interactive_metrics():
    names = list(MAPE_val.columns)
    # colors, if less than defined, create evenly spaced colors
    colors = ['olive', 'goldenrod', 'darkred', 'grey', 'midnightblue', 'crimson',
    'enta', 'darkslategray', 'darkslategrey', 'darkturquoise', 'darkviolet', 'deeppink',
    'ay', 'dodgerblue', 'firebrick', 'floralwhite', 'forestgreen', 'fuchsia', 'gainsboro',
    'goldenrod', 'gray', 'green', 'greenyellow', 'honeydew', 'hotpink', 'indianred',
    'khaki', 'lavender', 'lavenderblush', 'lawngreen', 'lemonchiffon', 'lightblue',
    'n']
    if len(names) > len(colors):
        evenly_spaced_interval = np.linspace(0, 1, len(names))
        colors = [cm.jet(x) for x in evenly_spaced_interval]

    # plotly setup
    fig = go.Figure()

    # Add one ore more traces
    for indx in range(len(names)):
        fig.add_traces(go.Line(y=MAPE_test[names[indx]],
                               name=names[indx],
                               fillcolor='blue',
                               mode='lines',
                               line=dict(color=colors[indx])))

    # construct menus
    updatemenus = [{['buttons': [ {'method': 'update',
                                    'label': 'MAPE Test',
                                    'args': [ {'y': [MAPE_test[names[indx]] for i in
                                         )])}, ]},
                    ],
                    },
                    {['method': 'update',
                      'label': 'MAPE Validation',
                      'args': [ {'y': [MAPE_val[names[indx]] for i in
                                         )])}, ]},
                    ],
                    },
                    {['method': 'update',
                      'label': 'RMSE Test',
                      'args': [ {'y': [RMSE_test[names[indx]] for i in
                                         )])}, ]},
                    ],
                    },
                    {['method': 'update',
                      'label': 'RMSE Validation',
                      'args': [ {'y': [RMSE_val[names[indx]] for i in
                                         )])}, ]},
                    ],
                    },
                    {['method': 'update',
                      'label': 'R2 Test',
                      'args': [ {'y': [R2_test[names[indx]] for i in
                                         )]}, ],
                      {['method': 'update',
                        'label': 'R2 Validation',
                        'args': [ {'y': [R2_val[names[indx]] for i in
                                         )]}]}]}]
```

```
)])},]],[  
    'direction':  
    'showactive'  
  
    # update layout with buttons.  
    # Code below is commented out  
    fig.update_layout(updatemenus=
```

Interactive Visualization of predictions

Execution of the code below will start an interactive plot (as demonstrated below) in the default browser. On the enable the user to choose of which model for which dataset and for which timestep the predictions shall be shown plot the kernel needs to be interrupted.

```
! bokeh serve --show visualization.py
```

The figure is a line graph titled "Predictions vs Actual Values for Test Set on Step 10 calculated by Naive Shift Model". The y-axis is labeled "Power Loss (normed)" and ranges from 0 to 0.7. The x-axis is labeled "Time" and shows dates from 4/22 to 4/28. Two lines are plotted: a blue line for "Predictions" and a red line for "Actual Values". Both lines show a sharp increase in power loss starting around April 22nd, peaking between 0.65 and 0.7 on April 23rd and 24th, and then gradually decreasing towards the end of the period. A sidebar on the right side of the plot area contains icons for "Model", "Data", "Step", and other settings.

```
MAE_test = get_erros(me
MAPE_val = get_erros(met
metric ='RMSE'
RMSE_test = get_erros(me
RMSE_val = get_erros(met
metric ='R2'
R2_test = get_erros(metr
R2_val = get_erros(metr
```

For a first comparison of error metrics on the validation and test set, the code below executes an interactive plot (as demonstrated below) within the current notebook.

```
plot_interactive_metrics()
```

The chart displays the Mean Absolute Percentage Error (MAPE) for seven different models over 18 time steps. The y-axis ranges from 0.1 to 0.6. The models are: fbprophet (green), moving average model (orange), multivariate LstmNN (red), multivariate LstmNN PeepHole (grey), naive model (dark blue), univariate LstmNN (pink), and univariate LstmNN PeepHole (light blue). The multivariate models generally show higher MAPE values compared to their univariate counterparts.

```
#plot_interactive_metrics()
```

Performance Analysis using TensorBoard

TensorBoard can be used to analyse the performance of different LSTMs over the traing epochs. The TensorBoard can be started via the following command

```
!tensorboard --logdir logs/fit
```

A Screenshot of the TensorBoard below. Within TensorBoard, the LSTMs are organized by the model name and the prediction step.

The TensorBoard interface shows a line graph of the epoch mean absolute percentage error (MAPE) over 28 training steps. Multiple colored lines represent different LSTM configurations, all showing a general downward trend as the number of steps increases. The legend at the bottom lists the models and their corresponding colors:

Name	Value	Step	Time	Relative
20201123-2331_2layer_50neurons_00Dropout_Peephole_s1/train	21.93	13	23.11.2020, 23:39:53	8m 15s
20201123-2331_3layer_15neurons_00Dropout_Peephole_s1/train	23.85	13	23.11.2020, 23:43:19	11m 19s
20201123-2331_3layer_30neurons_00Dropout_Peephole_s1/train	22.97	13	23.11.2020, 23:44:13	11m 52s
20201123-2331_3layer_50neurons_00Dropout_Peephole_s1/train	21.81	13	23.11.2020, 23:46:22	13m 28s
20201123-2331_4layer_20neurons_00Dropout_Peephole_s1/train	23.38	13	23.11.2020, 23:50:33	16m 2s

#!tensorboard --logdir logs/fit

Static Plots incl. Error Tables for deeper Investigation

For an indepth analysis, the error metrics can be inspected using the follwoing function.

```
# MAPE in Test Set
plot_static_metrics(MAPE_test.T, metric='MAPE', val_or_test='test')
MAPE_test.T
```

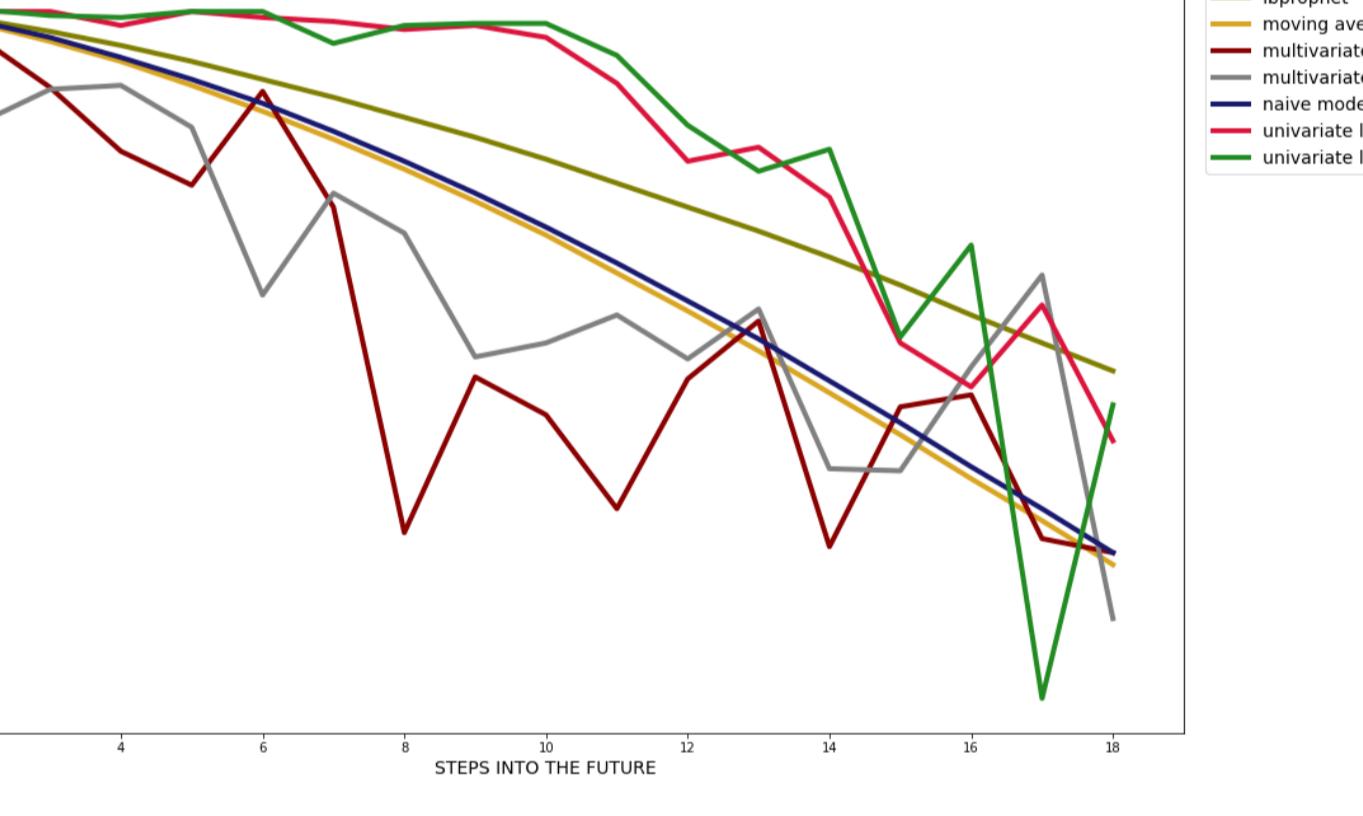
Step	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18
fbprophet	0.112	0.154	0.186	0.223	0.254	0.283	0.314	0.346	0.373	0.397	0.418	0.438	0.460	0.484	0.506	0.527	0.545	0.562
moving average model	0.124	0.157	0.197	0.233	0.267	0.301	0.336	0.370	0.399	0.426	0.451	0.473	0.499	0.525	0.546	0.567	0.587	0.604
multivariate IstmNN	0.359	0.366	0.415	0.451	0.480	0.441	0.467	0.528	0.462	0.494	0.526	0.619	0.514	0.516	0.495	0.554	0.546	0.602
multivariate IstmNN PeepHole	0.372	0.464	0.415	0.442	0.444	0.512	0.433	0.507	0.522	0.497	0.518	0.483	0.464	0.539	0.591	0.522	0.509	0.645
naive model	0.112	0.155	0.190	0.230	0.262	0.296	0.329	0.366	0.395	0.422	0.448	0.470	0.494	0.520	0.543	0.565	0.585	0.604
univariate IstmNN	0.127	0.120	0.171	0.289	0.146	0.186	0.217	0.267	0.195	0.217	0.304	0.458	0.349	0.392	0.389	0.524	0.507	0.515
univariate IstmNN PeepHole	0.163	0.115	0.165	0.155	0.159	0.158	0.190	0.183	0.254	0.345	0.283	0.279	0.330	0.411	0.406	0.403	0.411	0.413

MAPE for each predicted timestep into the future (test set)

The chart displays the Mean Absolute Percentage Error (MAPE) for seven different time series forecasting models over 18 predicted timesteps. The models are color-coded as follows: fbprophet (green), moving average model (blue), multivariate IstmNN (orange), multivariate IstmNN PeepHole (red), naive model (purple), univariate IstmNN (yellow), and univariate IstmNN PeepHole (pink). The y-axis represents the MAPE, ranging from 0.4 to 0.6. The x-axis represents the timestep, ranging from 1 to 18. The chart shows that most models show an increasing trend in MAPE over time, with the multivariate IstmNN PeepHole model generally having the highest MAPE values, peaking around 0.62 at timestep 18.

The chart displays the coefficient of determination (R2) for seven different models as they predict up to 18 steps into the future. The y-axis is logarithmic, ranging from 0.1 to 0.2. The x-axis represents the 'STEPS INTO THE FUTURE' from 2 to 18. The legend identifies the models: fbprophet (blue), moving average model (orange), multivariate IstmNN (red), multivariate IstmNN PeepHole (green), naive model (purple), univariate IstmNN (pink), and univariate IstmNN PeepHole (yellow). Most models show a general decline in R2 as the prediction horizon increases, with the univariate IstmNN PeepHole model showing a significant dip around step 17.

Step	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18
fbprophet	0.112	0.154	0.186	0.223	0.254	0.283	0.314	0.346	0.373	0.397	0.418	0.438	0.460	0.484	0.506	0.527	0.545	0.562
moving average model	0.124	0.157	0.197	0.233	0.267	0.301	0.336	0.370	0.399	0.426	0.451	0.473	0.499	0.525	0.546	0.567	0.587	0.604
multivariate IstmNN	0.359	0.366	0.415	0.451	0.480	0.441	0.467	0.528	0.462	0.494	0.526	0.619	0.514	0.516	0.495	0.554	0.546	0.602
multivariate IstmNN PeepHole	0.372	0.464	0.415	0.442	0.444	0.512	0.433	0.507	0.522	0.497	0.518	0.483	0.464	0.539	0.591	0.522	0.509	0.645
naive model	0.112	0.155	0.190	0.230	0.262	0.296	0.329	0.366	0.395	0.422	0.448	0.470	0.494	0.520	0.543	0.565	0.585	0.604
univariate IstmNN	0.127	0.120	0.171	0.289	0.146	0.186	0.217	0.267	0.195	0.217	0.304	0.458	0.349	0.392	0.389	0.524	0.507	0.515
univariate IstmNN PeepHole	0.163	0.115	0.165	0.155	0.159	0.158	0.190	0.183	0.254	0.345	0.283	0.279	0.330	0.411	0.406	0.403	0.411	0.413



This chart provides a detailed view of the R2 values for each individual timestep from 2 to 18. The y-axis ranges from 0.65 to 1.00. The x-axis is labeled 'STEPS INTO THE FUTURE'. The legend includes all seven models. The univariate IstmNN PeepHole model shows the most volatility, with a sharp dip to approximately 0.65 at step 17. The multivariate IstmNN model starts with the highest R2 value (around 0.98) but shows a steady decline after step 6. The fbprophet model maintains a relatively high R2 value (around 0.95-0.97) across most steps.

Step	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18
fbprophet	0.98	0.97	0.97	0.96	0.95	0.94	0.93	0.92	0.91	0.90	0.89	0.88	0.87	0.86	0.85	0.84	0.83	0.82
moving average model	0.98	0.97	0.97	0.96	0.95	0.94	0.93	0.92	0.91	0.90	0.89	0.88	0.87	0.86	0.85	0.84	0.83	0.82
multivariate IstmNN	0.98	0.97	0.97	0.96	0.95	0.94	0.93	0.92	0.91	0.90	0.89	0.88	0.87	0.86	0.85	0.84	0.83	0.82
multivariate IstmNN PeepHole	0.98	0.97	0.97	0.96	0.95	0.94	0.93	0.92	0.91	0.90	0.89	0.88	0.87	0.86	0.85	0.84	0.83	0.82
naive model	0.98	0.97	0.97	0.96	0.95	0.94	0.93	0.92	0.91	0.90	0.89	0.88	0.87	0.86	0.85	0.84	0.83	0.82
univariate IstmNN	0.98	0.97	0.97	0.96	0.95	0.94	0.93	0.92	0.91	0.90	0.89	0.88	0.87	0.86	0.85	0.84	0.83	0.82
univariate IstmNN PeepHole	0.98	0.97	0.97	0.96	0.95	0.94	0.93	0.92	0.91	0.90	0.89	0.88	0.87	0.86	0.85	0.84	0.83	0.82

The chart displays the Root Mean Square Error (RMSE) for three different model configurations over 10 timesteps into the future. The y-axis represents RMSE, ranging from 0.00 to 0.10. The x-axis represents the timestep index from 1 to 10. The legend identifies the series:

- univariate LSTMNN**: Represented by a blue line.
- univariate LSTMNN PeepHole**: Represented by a red line.

Timestep	univariate LSTMNN	univariate LSTMNN PeepHole
1	0.127	0.163
2	0.120	0.115
3	0.171	0.165
4	0.289	0.155
5	0.146	0.159
6	0.186	0.158
7	0.217	0.190
8	0.267	0.183

The figure is a line graph titled "RMSE vs STEPS INTO THE FUTURE". The y-axis is labeled "RMSE" and ranges from 0.00 to 0.08. The x-axis is labeled "STEPS INTO THE FUTURE" and ranges from 0 to 18. There are four data series:

- univariate LSTMNN** (red line): Starts at ~0.02, peaks at ~0.08 around step 8, and ends at ~0.075.
- univariate LSTMNN Peep Hole** (green line): Remains low until step 6, then rises steadily to ~0.08 by step 18.
- multivariate LSTMNN** (blue line): Starts at ~0.02, stays below the red line until step 10, and then rises to ~0.08 by step 18.
- multivariate LSTMNN Peep Hole** (grey line): Starts at ~0.025, peaks at ~0.07 around step 9, and ends at ~0.065.

 The legend is located in the top right corner of the plot area.

```
# MAE in Test Set
plot_static_metrics(MAE_test.T, met
MAPE_test.T
```



2 4 6 8 10 12 14 16 18

STEPS INTO THE FUTURE

Use Case - Power to Gas

The data in the capstone project comes from the windfarm [Twistriegen](#) in the south-east of Bremen. The windfarm consists of 13 wind turbines with a total installed capacity of **22.8 MW**. With this knowledge the normed power loss values of the forecasts can be transformed into a simple hands-on use case.

Usecase Introduction: Using potentially lost power

The diagram illustrates the hydrogen production and distribution process. It starts with wind turbines (represented by three blades) on the left. One path from the turbines goes through a 'Feed-In Mgmt Prediction' block (containing a sun and wind icon) and then a 'DEMAND SIDE MANAGEMENT SCHEDULING' block (containing a green switch icon). This path leads to an 'Electrolyzer' (represented by vertical bars) and then to a 'Hydrogen Tank' (represented by a grey cylinder). The hydrogen tank is connected to a 'Fuel Cell Bus' (represented by a black bus with 'FUEL CELL BUS' written on it). Another path from the turbines goes directly to a 'Feed-In Management order by Transmission System Operator' block (containing a red switch icon). This block also controls the 'Electrolyzer'. In the bottom right corner, there is a logo for 'tennet'.

Assumptions / Simplifications:

- A 100% efficient [electrolyser](#) requires 39 kWh to produce 1 kg of hydrogen. The devices today require as much as **48 kWh/kg**.
- A [fuel cell bus](#) with a pressure tank containing around **40 kg** of hydrogen (storage pressure 350 bar) can be refuelled in around seven minutes. With one tank filling, the bus can travel about **300 km** - 350 km.
- Ramp up times, Supply Chain Problems such as logistics and Storage Limitations are not taken into account.

Approach:

- To keep the example simple only the best-performing model will be evaluated, which is the **univariate LSTM-Model with**

- It is calculated how much otherwise lost power could have been used for hydrogen production, if the forecast would have been done **10 minutes**, **1.5 hours** and **3 hours** ahead for the 10-day test interval from **20/04/2019 06:00** to **29/04/2019 06:00**.
- The results are converted into the mileage that a fuel cell bus could otherwise have achieved.

```

df_uni_lstm_peephole_test_test["y_test_observed Step 9"],
df_uni_lstm_peephole_test_test["y_test_observed Step 18"]], axis=1)

# grabbing the time series that need to be integrated
y1_forecasted = combined_df[['y_test_pred Step 1','y_test_observed Step 1']].min(axis=1)
y9_forecasted = combined_df[['y_test_pred Step 9','y_test_observed Step 9']].min(axis=1)
y18_forecasted = combined_df[['y_test_pred Step 18','y_test_observed Step 18']].min(axis=1)
y1_actual = combined_df['y_test_observed Step 1']
y9_actual = combined_df['y_test_observed Step 9']
y18_actual = combined_df['y_test_observed Step 18']

# calculating the power loss that could have been predicted for the 10-day test set, if forecasts
# were shifted 10 days into the future

```

```

accumulated_predicted_loss9 = round((22.8 * simps(y9_forecasted, dx=1/6)),1)
accumulated_predicted_loss18 = round((22.8 * simps(y18_forecasted, dx=1/6)),1)
accumulated_predicted_loss = [accumulated_predicted_loss1, accumulated_predicted_loss9, accumulated_predicted_loss18]

# calculating the actual power loss for the 10-day test set
accumulated_actual_loss1 = int((22.8 * simps(y1_actual, dx=1/6)))
accumulated_actual_loss9 = int((22.8 * simps(y9_actual, dx=1/6)))
accumulated_actual_loss18 = int((22.8 * simps(y18_actual, dx=1/6)))
accumulated_actual_loss = [accumulated_actual_loss1, accumulated_actual_loss9, accumulated_actual_loss18]

# calculating the percentage of power loss that could have been predicted
percentage_predicted1 = round((accumulated_predicted_loss1 / accumulated_actual_loss1)*100,1)
percentage_predicted9 = round((accumulated_predicted_loss9 / accumulated_actual_loss9)*100,1)
percentage_predicted18 = round((accumulated_predicted_loss18 / accumulated_actual_loss18)*100,1)
percentage_predicted = [percentage_predicted1,percentage_predicted9,percentage_predicted18]

# converting the saved power into the mileage that a fuel cell bus could otherwise have achieved
buskilometers1 = int((accumulated_predicted_loss1 / (48/1000) * (300/40)))
buskilometers9 = int((accumulated_predicted_loss9 / (48/1000) * (300/40)))

```

```

data_tuples = list(zip(accumulated_actual_loss, accumulated_predicted_loss, percentage_predicted, busklioneters))
results = pd.DataFrame(data_tuples, columns=['Lost Power in MWh', 'Lost Power (forecasted) in MWh',
                                              'Percentage in %', 'Potential Mileage of Fuel Cell Bus in km'],
                        index = ["Prediction 10 minutes ahead", "Prediction 1.5 hours ahead", "Prediction 3 hours ahead"])
results

```

	Lost Power in MWh	Lost Power (forecasted) in MWh	Percentage in %	Potential Mileage of Fuel Cell Bus in km
Prediction 10 minutes ahead	503	495.0	98.4	77343
Prediction 1.5 hours ahead	503	451.9	89.8	70609
Prediction 3 hours ahead	503	322.3	64.1	50359

Results: 3 hour ahead forecast for hydrogen production

- The short term forecasts might not be realistically applicable for the scheduling of hydrogen production due to logistics and other concerns, but even for a 3 hours ahead forecast the accumulated power for over 50.000 km of potential fuel cell bus mileage could have been saved over the short amount of 10 days.

Future Work

APIs for live predictions

- [climacell Weather Data](#)

```
u.encoding = 'latin1'
df = u.load()

#### DATA SELECTION ####
start_index = pd.to_datetime("2019-04-26 00:00:00 +00:00")
end_index = pd.to_datetime("2019-05-12 00:00:00 +00:00")
f_to_plot = ['elevation', 'azimuth']

#reduced timeframe
df_reduced = df[(df.index <= end_index) & (df.index >= start_index)]

#### PLOTTING #####
color=iter(['olive', 'goldenrod', 'darkred', 'sandybrown', 'sienna', 'grey'])
```

The figure consists of two vertically stacked line plots. The top plot is titled "elevation" and shows a signal with a periodic oscillation between approximately -10 and 50. The bottom plot is titled "azimuth" and shows a signal with a periodic oscillation between approximately 180 and 250. Both plots have x-axis labels indicating time steps from 0 to 100.

```
for f, axes in zip(f_co_ploc, axs):
    c=next(color, 'goldenrod')
    axes.plot(df_reduced[f], color=c, lw=3)
    axes.set_title(f, fontsize=14)
    axes.set_xlim(left=start_index, right=end_index)
fig.tight_layout()
plt.show(fig);
```

Forecast of GFS Data for Feed-In Mgmt predictions >1 timestep

This is the final project at the Data Science Bootcamp [@neuefische](#). This project has been developed in the fall of 2020 by Tjade Appel ([LinkedIn](#) / [GitHub](#)) and Jonas Jaenicke ([LinkedIn](#) / [GitHub](#)). Please feel free to contact us.